

Analysis plan for AlgorithmWatch TAMI Dutch politics data

Boris Reuderink
boris@cortext.nl

February 26, 2021

Contents

1 Introduction

Towards a Monitoring of Instagram (TAMI) is a controlled experiment that will shed light on how items are chosen by Instagram’s algorithm, focusing on political parties in the Netherlands. The results will be shared with NOS to produce a series of articles. Then project will be successful when Instagram users in the Netherlands know more about our findings (measured by engagement on NOS and their Instagram feed) and if Instagram reacts meaningfully to our findings.

1.1 Aims

The aim of the data analysis is to reveal if specific types of posts on Instagram are overrepresented (i.e. are relatively more likely to show up in a users feed). We distinguish between two types of overrepresentation: 1) overrepresentation driven by post popularity, such as for viral content, and 2) overrepresentation driven by biases in Instagram’s recommendations. We are particularly interested in the latter type of overrepresentation. Note that this biased overrepresentation might differ from user to user. With the analysis described below, we collect evidence to answer the following research questions:

RQ1: What posts are overrepresented in the encounters? To answer this question, we need to describe what posts of Dutch politicians on Instagram are overrepresented in general. This might be due to popularity of the post, or due to biases in Instagram’s recommendations. We are interested in topics, as indicated by hashtags and by GoogleVision results for the post’s images.

RQ2: What posts are overrepresented given their popularity? Since overrepresentation is likely to be driven by popularity, we analyse what types of posts are overrepresented if we correct for effects of post popularity. Remaining variations in post popularity can indicate biases in Instagram’s recommendations.

RQ3: Is the overrepresentation the same in all subgroups? There might be biases that promote user engagement, such as bias for controversial content. But we might also find biases due to personalization. Biases due to personalization are expected to vary from user to user, while consistent biases are present for all users. We test the consistency of biases by evaluating the predictive performance of bias models on users that were not used for model development.

RQ4: Who benefit from these discrepancies? Finally, we can estimate per politician how much their posts were benefitting from a consistent algorithm bias of Instagram’s recommendations.

1.2 Data generation

The data from the TAMI experiment was generated as follows. A list with accounts from Dutch politicians was used to monitor all posts by these accounts in a certain period. A group of voluntary data donors ran a browser-plugin that retrieved the top posts in the timeline at hourly intervals, and records if a post from the monitored accounts was encountered. Donors were asked to follow a particular set of three politicians.

We have no information on the exact ranking algorithm that Instagram uses to determine what posts to present in the timeline of a user. Based on general properties of recommendation systems we expect that the probability of encountering a post is based on: the popularity of the post, the time of day, and day of week, the content of the post, and content preferences of the user. A graph of the assumed data generating processes is shown in Figure ??.

2 Analysis

First we pre-process data using database queries into a representation that facilitates easy analysis. We select time periods where a donor could have observed a particular post, and model the probability of an encounter in this period by fitting a logistic regression model. Properties of this model can answer RQ1. Next, we add features (exogenous variables) for post popularity. Properties of this model can answer RQ2. To evaluate the consistency of the found biases, we evaluate the properties of the found model on held-out donors. Due to the small number of donors we use cross-validation, where the procedure is repeated for different sets of held-out donors. Finally, we make take the consistent effects to estimate the expected algorithm bias per account, and per political party (RQ4).

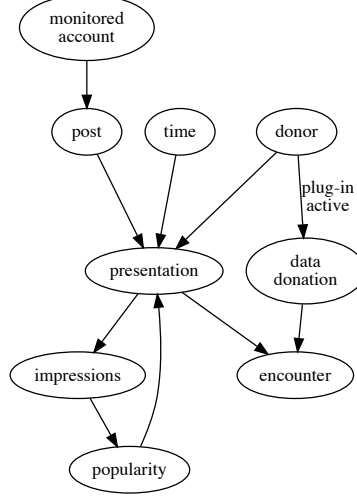


Figure 1: Schematic overview of the data generating process. Note that there is a feedback loop where increased impressions lead to increased popularity, which in turn leads to increased presentation of a post.

2.1 Preprocessing

We organize the analysis around potential encounters of posts of monitored politicians by donors at hourly intervals. Only the first 7 days after post-creation are considered. For each of potential encounter, we store the following dynamic properties at hourly intervals: the number of likes, the number of comments, and the presence of an encounter. This data is stored in a table T_d .

We use information from posts to model the probability of an encounter given a certain post. To this end, we prepare a representation of posts with properties that can explain the probability of an encounter. It is based on results from the Google Vision API (safe search, label annotations), hashtags from the post text, and the number of followers of the politician’s account. This static data is stored in table T_s .

2.2 Modeling

In the previous section we introduced the tables with data that does not vary with time (T_s) and with data that does vary with time (T_d). In this section we outline a series of analysis steps that will lead to evidence to answer the research questions introduced in section ??.

2.2.1 Post overrepresentation

RQ1 is about what types of posts are overrepresented in general. Since we organized our data around potential encounters, we can use the static properties from T_s to predict the presence of an encounter in the time window with a logistic-regression model. In addition, we include a donor ID in the predictors to automatically account for difference in the amount of posts from other followed accounts. Coefficients of this regression model indicate how different post properties contribute to overrepresentation. Alternatively, we can use feature permutation to quantify how much a specific post property contributes to predictive performance.

2.2.2 Bias

For RQ2 we need to know what types of post are overrepresented given their popularity. We add information on post popularity and repeat the analysis in section ??.

2.2.3 Consistency of bias

Effects of overrepresentation found while controlling for post popularity indicate a bias of Instagram's recommendations. This bias might be consistent, but it might also be a consequence of personalization. To find consistent biases, we evaluate how much individual post properties contribute to predicting encounters on held-out donors. If a post property consistently contributes to accurate predictions of engagements for other donors we have found a consistent bias. This is quantified using permutation feature importance on held-out data.

Due to the low-number of donors, we have to train multiple models, where different donors are used as held-out data. For each of these models, we have to perform the analysis, and combine the results to construct a final set of biases.

2.2.4 Interpreting bias at account and party level

Finally, the biases found in section ?? can be estimated on the user account level, and aggregated on the political party level. This can be done by setting the models coefficients for popularity features to zero, making a prediction, and performing the aggregation.